

[jiahua-zou](#) / [lending-club-loan-project](#) Private

Branch: master ▼

[lending-club-loan-project](#) / [logit_svm_rf.ipynb](#)

Find file

Copy path

 **JiahuaZou** add visualization, change file name

1a9bda5 7 days ago

[0 contributors](#)

1522 lines (1521 sloc) 681 KB

```
In [1]: library(dplyr)
library(tidy)
library(ISLR)
library(MASS)
library(ggplot2)
library(e1071)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Attaching package: 'MASS'

The following object is masked from 'package:dplyr':

select

```
In [2]: df <- read.csv('cleaned_loan_data.csv')
head(df$loan_rate, 20)
```

```
1. 2
2. 2
3. 2
4. 1
5. 2
6. 2
7. 2
8. 2
9. 2
10. 1
11. 2
12. 1
13. 2
14. 2
15. 1
16. 0
17. 0
18. 1
19. 2
20. 1
```

```
In [3]: df2 <- df %>% filter(loan_rate != 2) %>% dplyr::select(-c(policy_code, id, X, member_id, term,
out_prncp_inv)) %>% select_if(is.numeric)
```

```
df2$loan_rate <- as.factor(df2$loan_rate)
head(df2$loan_rate)
```

```
1. 1
2. 1
3. 1
4. 1
5. 0
6. 0
```

```
In [4]: useful_var = c("loan_amnt", "funded_amnt", "funded_amnt_inv",
                       "int_rate", "installment", "dti", "delinq_2yrs", "inq_last_6mths",
                       "open_acc", "pub_rec", "revol_bal", "revol_util", "total_acc",
                       "total_rec_int", "total_rec_late_fee", "collections_12_mths_ex_med",
                       "acc_now_delinq", "tot_coll_amt", "tot_cur_bal", "total_rev_hi_lim", "loan_rate")
```

```
In [5]: df3 <- df2[, useful_var]
colnames(df3)
```

```
1. 'loan_amnt'
2. 'funded_amnt'
3. 'funded_amnt_inv'
4. 'int_rate'
5. 'installment'
6. 'annual_inc'
7. 'dti'
8. 'delinq_2yrs'
9. 'inq_last_6mths'
10. 'open_acc'
11. 'pub_rec'
12. 'revol_bal'
13. 'revol_util'
14. 'total_acc'
15. 'out_prncp'
16. 'total_pymnt'
17. 'total_pymnt_inv'
18. 'total_rec_prncp'
19. 'total_rec_int'
20. 'total_rec_late_fee'
21. 'recoveries'
22. 'collection_recovery_fee'
23. 'last_pymnt_amnt'
24. 'collections_12_mths_ex_med'
25. 'acc_now_delinq'
26. 'tot_coll_amt'
27. 'tot_cur_bal'
28. 'total_rev_hi_lim'
29. 'loan_rate'
```

```
In [6]: set.seed(1)
train=sample(nrow(df3), round(0.75 * nrow(df3)))
trainset=df3[train,]
testset=df3[-train,]
```

```
In [7]: glm.fit2 = glm(loan_rate~., data = trainset, family = binomial)
summary(glm.fit2)
```

Call:

```
glm(formula = loan_rate ~ ., family = binomial, data = trainset)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-4.0322	-0.9056	0.5950	0.7818	4.7886

```

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)    3.367e+00  3.678e-02  91.537 < 2e-16 ***
loan_amnt      -3.203e-04  3.972e-04  -0.806  0.42004
funded_amnt     1.454e-03  4.605e-04   3.157  0.00159 **
funded_amnt_inv -1.215e-03  2.312e-04  -5.256  1.47e-07 ***
int_rate       -1.078e-01  1.663e-03 -64.820 < 2e-16 ***
installment     1.778e-03  8.138e-05  21.850 < 2e-16 ***
dti            -3.396e-02  8.251e-04 -41.159 < 2e-16 ***
delinq_2yrs    -1.020e-01  7.090e-03 -14.380 < 2e-16 ***
inq_last_6mths -2.677e-02  5.954e-03  -4.496  6.92e-06 ***
open_acc       -3.197e-02  1.756e-03 -18.207 < 2e-16 ***
pub_rec        -5.060e-02  1.222e-02  -4.141  3.46e-05 ***
revol_bal      -6.037e-06  9.567e-07  -6.310  2.79e-10 ***
revol_util     -2.765e-03  3.755e-04  -7.364  1.78e-13 ***
total_acc      1.863e-02  7.478e-04  24.914 < 2e-16 ***
total_rec_int   8.251e-05  3.943e-06  20.927 < 2e-16 ***
total_rec_late_fee -8.738e-02  1.705e-03 -51.263 < 2e-16 ***
collections_12_mths_ex_med -3.251e-01  5.238e-02 -6.207  5.41e-10 ***
acc_now_delinq -1.620e-01  7.993e-02  -2.027  0.04267 *
tot_coll_amt    5.572e-08  4.133e-07   0.135  0.89276
tot_cur_bal     1.366e-06  5.506e-08  24.816 < 2e-16 ***
total_rev_hi_lim  6.846e-06  6.705e-07  10.210 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 183474 on 157733 degrees of freedom
Residual deviance: 164416 on 157713 degrees of freedom
AIC: 164458

```

Number of Fisher Scoring iterations: 5

```

In [8]: glm.probs=predict(glm.fit2, testset, type="response")
        glm.pred=rep(0, nrow(testset))
        glm.pred[glm.probs>0.5]=1
        table(glm.pred, testset$loan_rate)
        1-mean(glm.pred==testset$loan_rate)

```

```

glm.pred    0    1
0  2494  1510
1 11557 37017

```

0.248525999467458

```

In [8]: colnames(testset)

```

```


1. 'loan_amnt'
2. 'funded_amnt'
3. 'funded_amnt_inv'
4. 'int_rate'
5. 'installment'
6. 'dti'
7. 'delinq_2yrs'
8. 'inq_last_6mths'
9. 'open_acc'
10. 'pub_rec'
11. 'revol_bal'
12. 'revol_util'
13. 'total_acc'
14. 'total_rec_int'
15. 'total_rec_late_fee'
16. 'collections_12_mths_ex_med'
17. 'acc_now_delinq'
18. 'tot_coll_amt'
19. 'tot_cur_bal'
20. 'total_rev_hi_lim'

```

21. 'loan_rate'


```
In [9]: testset$loan_rate <- as.numeric(as.character(testset$loan_rate))
```

```
In [16]: fit = glm(loan_rate ~ int_rate, data=trainset, family=binomial)
newdat <- data.frame(int_rate=seq(min(testset$int_rate), max(testset$int_rate), len=100))
newdat$loan_rate = predict(fit, newdata=newdat, type="response")
plot(loan_rate~int_rate, data=testset, col="red4")
lines(loan_rate~int_rate, newdat, col="green4", lwd=2)
```



```
In [42]: ggplot(testset, aes(x=int_rate, y=loan_rate)) + geom_point() +
  geom_smooth(method = "glm", method.args = list(family = "binomial"), se = FALSE)
```

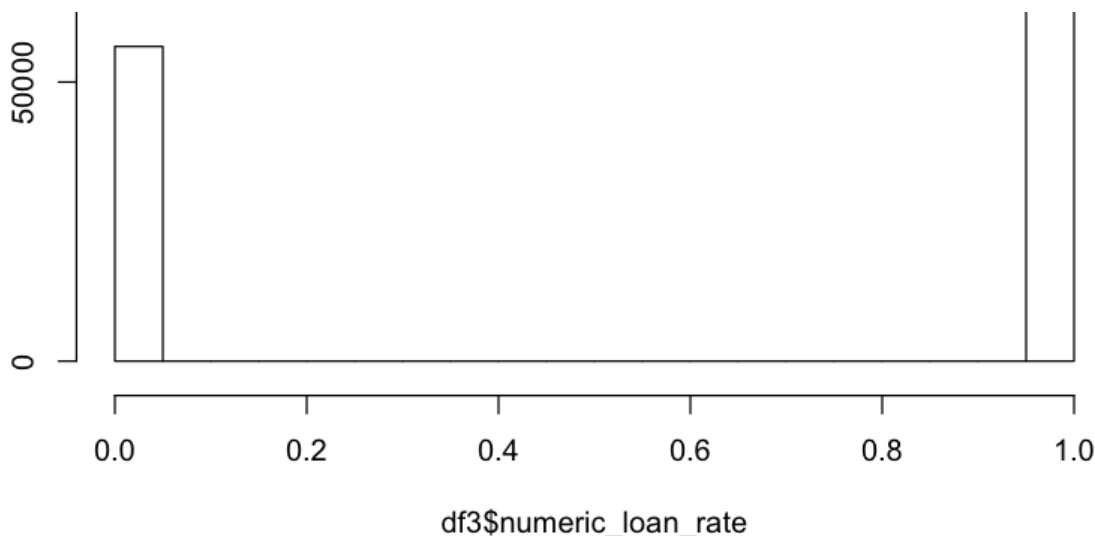




```
In [68]: hist(df3$numeric_loan_rate)
```

Histogram of df3\$numeric_loan_rate





```
In [17]: temp <- sample(nrow(df3), 30000)
df4 <- df3[temp, ]
```

```
In [18]: train = sample(nrow(df4), round(0.8*nrow(df4)))
length(train)

24000
```

```
In [19]: #set.seed(1)
#tune.out=tune(svm, loan_rate~., data=df4[train,], kernel = "sigmoid",
#ranges=list(cost=c(0.1, 1, 10, 100, 1000)))
#summary(tune.out)
```

```
In [30]: svmfit=svm(loan_rate~., data=df4[train,], kernel = "radial", cost = 0.1)
y_pred2 = predict(svmfit, newdata=df4[-train,])
```

```
In [31]: true=df4[-train, "loan_rate"]
table(true, y_pred2)
1-(sum(true == y_pred2)/nrow(df4[-train,]))

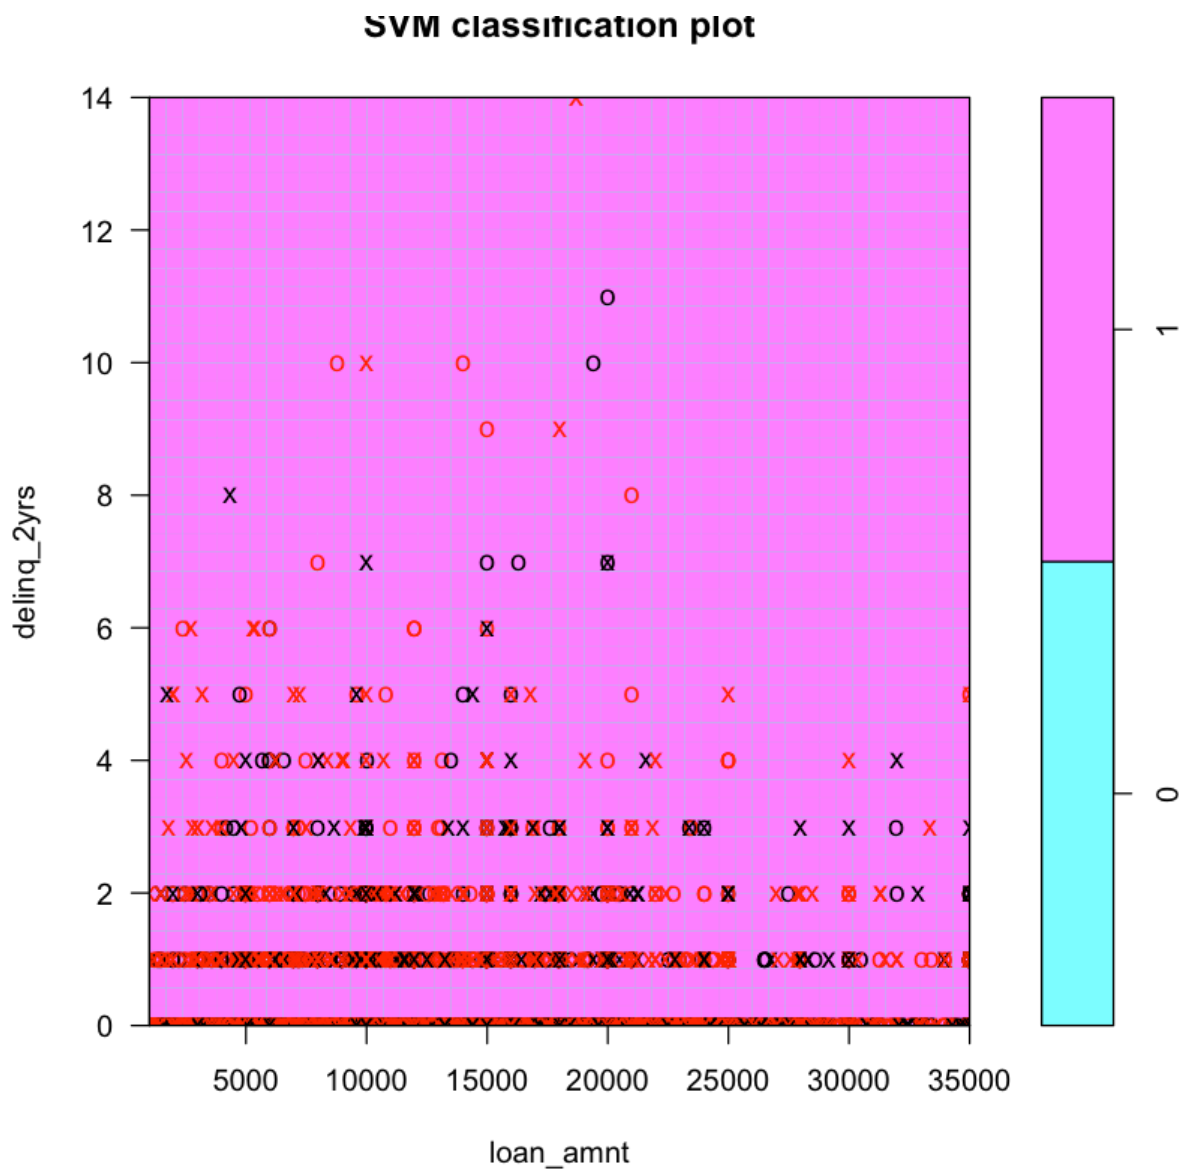
      y_pred2
true      0      1
  0    91 1504
  1    44 4361

0.258
```

```
In [32]: temp <- df4[-train,]
head(temp)
```

	loan_amnt	funded_amnt	funded_amnt_inv	int_rate	installment	dti	delinq_2yrs	inq_last_6mths	open_acc	pub
63706	8400	8400	8400	6.03	255.66	15.29	0	3	12	0
77293	8000	8000	8000	21.00	301.41	14.88	0	2	13	0
2773	10000	10000	10000	16.24	352.76	14.22	2	5	8	0
38063	14825	14825	14825	6.03	451.21	20.53	0	1	15	0
57878	28000	28000	28000	22.47	780.84	16.97	0	1	13	0
138046	10000	10000	10000	9.17	318.79	22.46	2	0	8	0

```
In [44]: plot(svmfit, temp, delinq_2yrs~loan_amnt)
```



```
In [45]: true2=df2[-train, "loan_rate"]
pred2=predict(svmfit, newdata=df2[-train,])
table(true2, pred2)
```

```
      pred2
true2    0      1
0    3815  46669
1     932 134896
```

```
In [46]: 1-(sum(true2 == pred2)/nrow(df2[-train,]))
```

```
0.255490789643179
```


```
In [52]: svmfit3 = svm(loan_rate~., data=df4[train,], kernel = "sigmoid", cost = 0.1)
y_pred3 = predict(svmfit3, newdata=df4[-train,])
true3=df2[-train, "loan_rate"]
pred3=predict(svmfit3, newdata=df2[-train,])
table(true3, pred3)
1-(sum(true3 == pred3)/nrow(df2[-train,]))
```

```
      pred3
true3    0      1
0   15881  34603
1   25384 110444
```

0.321970672849843

In [54]: $1 - ((15881 + 110444) / (15881 + 34603 + 25384 + 110444))$

0.321970672849843

In [49]: `plot(svmfit3, temp, delinq_2yrs~loan_amnt)`In [48]: `colnames(df4)`

1. 'loan_amnt'
2. 'funded_amnt'
3. 'funded_amnt_inv'
4. 'int_rate'
5. 'installment'
6. 'dti'
7. 'delinq_2yrs'
8. 'inq_last_6mths'
9. 'open_acc'


```

10. 'pub_rec'
11. 'revol_bal'
12. 'revol_util'
13. 'total_acc'
14. 'total_rec_int'
15. 'total_rec_late_fee'
16. 'collections_12_mths_ex_med'
17. 'acc_now_delinq'
18. 'tot_coll_amt'
19. 'tot_cur_bal'
20. 'total_rev_hi_lim'
21. 'loan_rate'

```

```
In [34]: train2 = sample(nrow(df3), round(0.8 * nrow(df3)))
length(train2)
```

```
168250
```

```
In [38]: trainset2 <- df3[train2,]
dim(trainset2)
```

```
1. 168250
2. 21
```

```
In [41]: testset2 <- df3[-train2, ]
dim(testset2)
```

```
1. 42062
2. 21
```

```
In [39]: library(randomForest)
rf_model <- randomForest(loan_rate~.,data = trainset2, keep.forest=TRUE)
```

```
In [40]: rf_model
```

```

Call:
randomForest(formula = loan_rate ~ ., data = trainset2, keep.forest = TRUE)
      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 4

      OOB estimate of  error rate: 24.03%
Confusion matrix:
      0      1 class.error
0 9856 35278 0.78162804
1 5157 117959 0.04188733

```


```
In [30]: varImpPlot(rf_model)
```

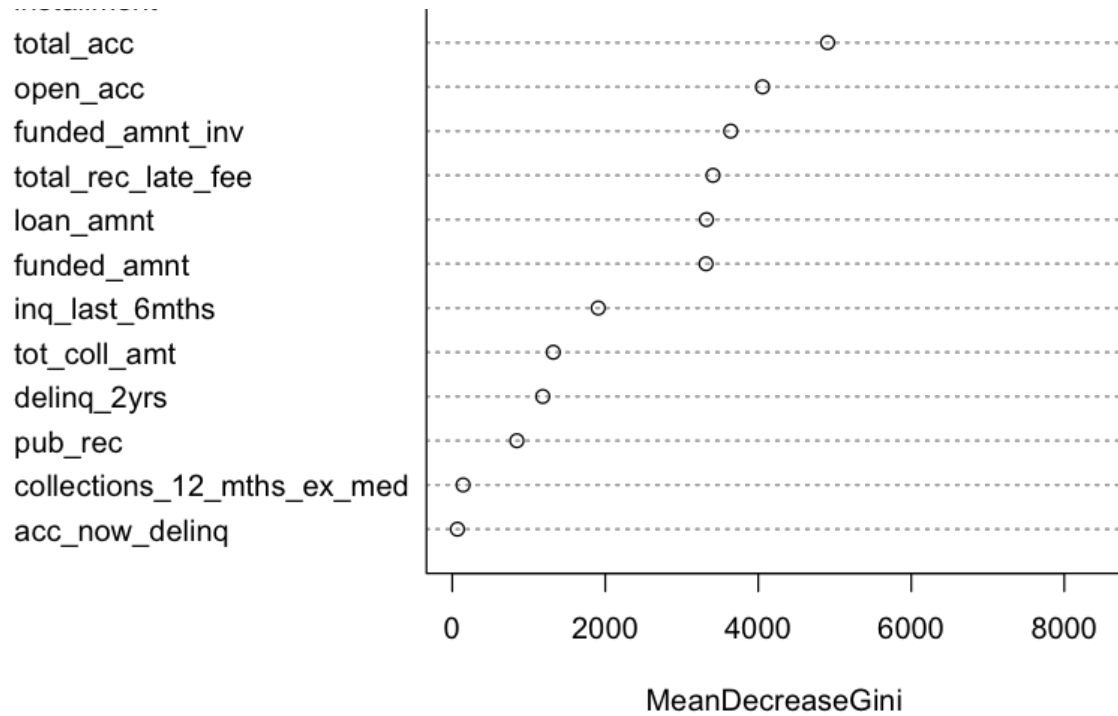
rf_model

```

int_rate
total_rec_int
dti
tot_cur_bal
revol_util
revol_bal
total_rev_hi_lim
installment

```






```
In [31]: importance(rf_model)
```

	MeanDecreaseGini
loan_amnt	3323.0490
funded_amnt	3317.6038
funded_amnt_inv	3640.2764
int_rate	8423.2912
installment	5588.9964
dti	7516.2359
delinq_2yrs	1183.2540
inq_last_6mths	1908.2250
open_acc	4056.0094
pub_rec	844.6215
revol_bal	5978.1019
revol_util	6251.2847
total_acc	4906.5944
total_rec_int	7733.7014
total_rec_late_fee	3408.2718
collections_12_mths_ex_med	142.0265
acc_now_delinq	67.7321
tot_coll_amt	1320.5947
tot_cur_bal	6805.3234
total_rev_hi_lim	5948.2921

```
In [43]: yhat_rf <- predict(rf_model, newdata = testset2)
plot(yhat_rf, testset2$loan_rate)
```

```
abline(0, 1)
```

```
Warning message in Ops.factor(yhat_rf, testset2$loan_rate):  
"-' not meaningful for factors"  
<NA>
```



```
In [44]: table(testset2$loan_rate, yhat_rf)  
1-(sum(testset2$loan_rate == yhat_rf)/nrow(testset2))
```

```
      yhat_rf  
      0      1  
0  2430  8812  
1  1145 29675  
0.236721981836337
```